

Floating Point


How machines see numbers

Prof. Nipun Batra

Mathematical Foundations for AI · IIT Gandhinagar

**A softmax computation that
returns nan**

617 steps of beautiful training, then `nan` forever:



Today we solve it — and derive the stable log-softmax computation used by major ML frameworks.

A decimal calculation worth checking

First, type this into any Python prompt:

```
>>> 0.1 + 0.2 == 0.3
False

>>> 0.1 + 0.2
0.30000000000000004
```

Not a Python bug. The same happens in C, Java, Rust, JavaScript, Excel, your calculator app...

Guess before we continue: is $0.5 + 0.25 == 0.75$ also `False`? Commit to an answer — we'll resolve it (and you'll see *exactly why*) within the hour.

Learning outcomes

By the end of this lecture you will be able to:

1. Explain why one fixed-point scale cannot cover a training run's full **dynamic range**, and how quantized ML works around that limit.
2. Decode and encode **IEEE-754 float32** by hand (sign / exponent / mantissa).
3. Reason about the **unevenly spaced** float number line: machine epsilon, $x + 1 == x$.
4. Predict **overflow / underflow**: why $\exp(89)$ breaks float32.
5. Spot **catastrophic cancellation** and rewrite formulas to dodge it.
6. Derive and apply the **log-sum-exp trick** and the numerically stable softmax.
7. Choose between **float32 / float16 / bfloat16** and explain why many accelerators favour bfloat16 or mixed precision.

Before floats: integers and fixed point

Integers · exact, but boxed in

A 32-bit integer stores any whole number in $[-2^{31}, 2^{31} - 1] \approx \pm 2.1 \times 10^9$ — **exactly**.

```
>>> 2**31 - 1          # largest int32
2147483647
```

- Arithmetic on integers is *perfect*: no rounding, ever.
- But there are no fractions — and beyond the box, NumPy `int32` wraps modulo 2^{32} . Languages differ: signed overflow in C is undefined, so never rely on wrapping unless the type specifies it.

Q. So how do we get decimals? First idea: just *fix* where the point goes.

Fixed point · a decimal point bolted in place

Split 32 bits: 16 for the whole part, 16 for the fraction (“16.16”):

- Smallest step: $2^{-16} \approx 0.000015$ — everywhere on the line
- Largest value: $2^{15} \approx 32,768$
- Example: $6.25 = 00000000000000110.0100000000000000_2$

Simple, fast, and used in DSP chips and retro game consoles. **But look at that range:** nothing above 33k, nothing (nonzero) below 0.000015 — and both limits are hard walls.

Why one fixed-point scale is not enough

One training run of one neural network routinely contains, **simultaneously**:

Quantity	Typical magnitude
probability of a rare token	10^{-30} and below
a small gradient	10^{-8}
a learning rate	10^{-3}
activations, logits	1 to 10^4
a loss spike, a softmax numerator	10^{10} and beyond

That's 40+ orders of magnitude in one program. One 32-bit fixed-point scale covers only a small window, so general-purpose training needs a format whose scale can move.

Quantized ML uses many local scales

Fixed point remains useful when the program does **not** rely on one global scale:

- assign a scale and zero-point to each tensor or channel,
- perform fast integer matrix products,
- rescale between layers using calibration data or learned quantization parameters.

Quantized inference works around fixed point's narrow range by using many locally chosen scales. General training still usually relies on floating point or mixed precision.

The rescue idea is 400 years old · scientific notation

Chemists never write Avogadro's number as 602, 214, 076, 000, 000, 000, 000. They write:

$$6.022 \times 10^{23}$$

- **A few significant digits** (the *mantissa*: 6.022) — the precision
- **A magnitude field** (the *exponent*: 23) — controls the range


Same digit budget describes atoms (10^{-27} kg) or stars (10^{30} kg): the point *floats* to where it's needed.

A floating-point number is scientific notation, in base 2, squeezed into 32 bits. That is the whole idea — the rest is bookkeeping.

IEEE-754: the 32 bits, dissected

The anatomy of a float32

float32 anatomy — the 32 bits of 6.25



$$\text{value} = (-1)^0 \times 1.1001_2 \times 2^{129 - 127} = 1.5625 \times 4 = 6.25$$

The decoding rule

$$\text{value} = (-1)^{\text{sign}} \times \underbrace{1.\text{mantissa}}_{24 \text{ significant bits}} \times 2^{\text{exponent} - 127}$$

- **Sign** (1 bit): 0 is positive, 1 is negative
- **Exponent** (8 bits): stored with a **bias of 127**, so the field [1, 254] covers powers 2^{-126} to 2^{127} — no separate sign needed for the exponent
- **Mantissa** (23 bits): binary normalized numbers always start “1.”, so the leading 1 is **not stored** — a free 24th bit of precision

Two exponent patterns are reserved: all-zeros (zero & *subnormals*, later ★★★) and all-ones (*inf* & *nan*, soon).

Step 1 — binary. $6.25 = 4 + 2 + 0.25 = 110.01_2$

Step 2 — normalize. $110.01_2 = 1.1001_2 \times 2^2$

Step 3 — sign. positive $\rightarrow s = 0$

Step 4 — exponent. $2 + 127 = 129 = 10000001_2$

Step 5 — mantissa. drop the leading “1.” $\rightarrow 100100000000000000000000$

0 10000001 100100000000000000000000 = 0x40C80000

Tutorial 1's worksheet has you do this by hand for -0.75 , 13.5 , and one nasty one...

Worked encoding · 0.1 — and the wheels come off

Multiply by 2, harvest the integer bit, repeat:

$$0.1 \xrightarrow{\times 2} \underline{0}.2 \xrightarrow{\times 2} \underline{0}.4 \xrightarrow{\times 2} \underline{0}.8 \xrightarrow{\times 2} \underline{1}.6 \xrightarrow{\times 2} \underline{1}.2 \xrightarrow{\times 2} \underline{0}.4 \dots$$

$$0.1 = 0.00011\overline{1}_2 = 0.00011001100110011\dots_2 \quad \text{— repeats forever!}$$

Like $\frac{1}{3} = 0.333\dots$ in decimal: a perfectly clean fraction, an **infinite** expansion in the wrong base. The mantissa holds 23 bits, so the machine must **round**:

```
>>> from decimal import Decimal
>>> Decimal(float(np.float32(0.1)))
Decimal('0.100000001490116119384765625')
```

There is **no 0.1 in floating point. There never was.**

Why `0.1 + 0.2 != 0.3`

What float64 *actually stores* (every digit below is exact):

You typed	The machine stored
<code>0.1</code>	0.1000000000000000000055511151231257827...
<code>0.2</code>	0.20000000000000000000111022302462515654...
<code>0.1 + 0.2</code>	0.30000000000000000000444089209850062616...
<code>0.3</code>	0.29999999999999999999888977697537484345...


The two calculations produce different adjacent floating-point values, so `==` returns `False`.

Do not use `==` when you mean “equal up to rounding” — for example, after inexact decimal arithmetic. Use `np.isclose(a, b) / math.isclose(a, b)`. Exact binary values, integer-valued floats, and sentinels can still be compared exactly.


The float number line

The number line, as the machine sees it

a toy float (3 mantissa bits): the gap doubles at every power of 2



real float32: spacing between consecutive numbers



Machine epsilon · the gap next to 1.0

Machine epsilon ε = the gap between 1.0 and the next representable float.

$$\varepsilon_{\text{float32}} = 2^{-23} \approx 1.19 \times 10^{-7} \quad \varepsilon_{\text{float64}} = 2^{-52} \approx 2.22 \times 10^{-16}$$

Anything smaller than about $\varepsilon/2$, added to 1, simply **vanishes**:

```
>>> np.float32(1.0) + np.float32(1e-8) == np.float32(1.0)
True                                     # 1e-8 fell into the gap

>>> 1.0 + 1e-8 == 1.0                  # float64's gap is much smaller
False
```

Rule of thumb: float32 carries **~7 reliable decimal digits**; float64 carries ~16.

At 2^{24} , adding 1 does nothing

The gap doubles at every power of 2. By $2^{24} = 16,777,216$ the gap between neighboring float32s is **2** — so +1 rounds right back:

```
>>> x = np.float32(16_777_216)      # 2**24
>>> x + np.float32(1) == x
True
>>> np.spacing(x)                   # the gap at x
2.0
```

A float32 counter that increments by 1 **stops counting at 16.7 million** — silently. Count in integers; accumulate long sums (losses, metrics) in float64.

Rounding · one tiny lie per operation

Each arithmetic result is computed exactly, then **rounded** to the nearest float:

$$\text{fl}(a \circ b) = (a \circ b)(1 + \delta), \quad |\delta| \leq \frac{\varepsilon}{2}$$

This standard relative-error model applies when the exact result is a finite normal number; underflow, overflow, and subnormals need separate treatment.

- Default mode: **round to nearest, ties to even** — ties go to the even last bit, so up/down rounds don't accumulate a bias
- Other IEEE modes exist (toward 0, toward $\pm\infty$) — used for interval arithmetic, not our concern

One operation = one tiny, *bounded* lie (10^{-8} relative in float32). The drama is never one lie — it's how lies **compound.**

When arithmetic escapes the range · `inf` and `nan`

IEEE-754 doesn't crash on impossible arithmetic — it returns special values:

Expression (NumPy)	Result	Meaning
<code>np.float64(1.0) / 0.0</code>	<code>inf</code>	division by zero / limiting value
<code>np.float64(-1.0) / 0.0</code>	<code>-inf</code>	
<code>np.float64(0.0) / 0.0</code>	<code>nan</code>	“not a number” — no defensible value
<code>inf - inf, inf / inf</code>	<code>nan</code>	
<code>np.log(-1.0), np.sqrt(-1.0)</code>	<code>nan</code>	

Most arithmetic with `nan` produces `nan`, so one invalid value can spread quickly through a computation graph. Comparisons are different: `nan == nan` is `False` (`x != x` is the classic test; prefer `np.isnan`).

Overflow arrives shockingly early

The range boundary isn't some astronomical corner case. For `exp`:

Computation	float32 (max $\approx 3.4 \times 10^{38}$)	float64 (max $\approx 1.8 \times 10^{308}$)
<code>exp(88)</code>	1.65×10^{38} ✓	fine
<code>exp(89)</code>	inf	fine
<code>exp(709)</code>	inf	8.2×10^{307} ✓
<code>exp(710)</code>	inf	inf
<code>exp(-104)</code>	0.0 (underflow)	fine

In float32, e^z overflows for z near 89; the opening logits around 1000 therefore overflow.

Catastrophic cancellation

Subtraction of near-equals is a digit shredder

Work in 8 significant decimal digits (like a float32-ish machine):

$$1.2345678 - 1.2345677 = 0.0000001$$

Both inputs had **8** correct digits. The result has **1** — and if the inputs were themselves rounded (they always are), that surviving digit is *pure noise*.

Adding same-sign numbers keeps relative error tame. Subtracting nearly equal numbers cancels the trustworthy leading digits and promotes the rounded-off garbage to the front.

Nothing “overflows” — the answer is just quietly wrong.

Worked example 1: the quadratic formula

Solve $x^2 - 10000x + 1 = 0$ in float32. The small root, two algebraically identical ways:

Naive (subtract near-equals: $b \approx \sqrt{b^2 - 4}$):


$$x = \frac{b - \sqrt{b^2 - 4}}{2} = \frac{10000 - 9999.9998...}{2} \xRightarrow{\text{float32}} \mathbf{0.0}$$

Stable (rationalize — same numbers, now *added*):

$$x = \frac{2}{b + \sqrt{b^2 - 4}} = \frac{2}{19999.9998} \xRightarrow{\text{float32}} \mathbf{1.0000 \times 10^{-4} \checkmark}$$

The naive formula returned **zero** — a 100% error — with no warning of any kind.

The two formulas, head to head



Same math on paper. In float32, one drifts to 100% error; the other sits at machine epsilon.

Worked example 2: two variance formulas

Textbooks give two equal formulas for variance. Try both on $[10000.0, 10000.1, 10000.2]$ (true variance ≈ 0.00667), in float32:

One-pass $\mathbb{E}[X^2] - (\mathbb{E}[X])^2$: two huge, nearly equal numbers —

$100001992.0 - 100002011.7 = -19.7$ (a negative variance!)

Two-pass $\mathbb{E}[(X - \bar{X})^2]$: subtract *first*, while numbers are small —

mean of $\{(-0.1)^2, 0^2, (0.1)^2\} = 0.00668 \checkmark$

Near 10^8 , float32's grid spacing is 8 — the true answer is a thousand times *smaller than the gap between representable numbers* there. Keep intermediate quantities small.

The ML fixes

Fix 1 · probabilities live in log-space

A language model scores a 1000-character text; each character has probability ≈ 0.03 :

$$P(\text{text}) = \prod_{i=1}^{1000} p_i \approx 0.03^{1000} \approx 10^{-1523} \Rightarrow \text{even float64 } \mathbf{0.0}$$

In log-space — products become sums, tiny becomes ordinary:

$$\log P(\text{text}) = \sum_{i=1}^{1000} \log p_i \approx 1000 \times (-3.51) = -3507 \quad \checkmark \text{ a boring, safe float}$$

Multiply probabilities $\rightarrow$ add log-probabilities. Many probabilistic objectives are negative log-likelihoods (L14); other losses need not be.

Fix 2 · log-sum-exp, derived in three lines

In log-space we still must *normalize*: compute $\log \sum_i e^{x_i}$ — but the e^{x_i} overflow! Pull out the max $m = \max_i x_i$:

$$\log \sum_i e^{x_i} = \log \sum_i e^m e^{x_i - m} = \log \left(e^m \sum_i e^{x_i - m} \right) = m + \log \sum_i e^{x_i - m}$$

Every exponent $x_i - m \leq 0$, so every $e^{x_i - m} \in (0, 1]$: **nothing can overflow** — and the largest term is exactly 1, so no total underflow either.

Check on the opening scores [1000, 1001, 1002]:

$$\text{LSE} = 1002 + \log(e^{-2} + e^{-1} + e^0) = 1002 + \log(1.503) = 1002.41 \quad \checkmark$$

Fix 3 · the stable softmax

Softmax is **shift-invariant** — multiply top and bottom by e^{-c} :



$$\text{softmax}(z)_i = \frac{e^{z_i}}{\sum_j e^{z_j}} = \frac{e^{z_i - c}}{\sum_j e^{z_j - c}} \quad \text{for any } c$$

Choose $c = \max_j z_j$. For $z = [1000, 1001, 1002]$:

	z_i	$z_i - 1002$	$e^{z_i - 1002}$	probability
naive	$e^{1000} = \text{inf}$		inf / inf	nan
stable		$-2, -1, 0$	0.135, 0.368, 1.0	0.090, 0.245, 0.665 ✓

Identical mathematics. One works on real hardware; the other burned down a training run.

The fix, in pictures



The fixes, in code

```
def softmax_naive(z):  
    e = np.exp(z)                # inf for z ~ 89+ → nan  
    return e / e.sum()  
  
def softmax_stable(z):  
    e = np.exp(z - z.max())      # largest exponent is now 0  
    return e / e.sum()          # same answer, can't overflow  
  
def logsumexp(x):  
    m = x.max()  
    return m + np.log(np.exp(x - m).sum())
```

```
>>> softmax_stable(np.array([1000., 1001., 1002.]))  
array([0.09003057, 0.24472847, 0.66524096])
```

Five lines you will reuse in Tutorial 1, in L14 (MLE), in L24 (cross-entropy), and in L26's language model.

Stable softmax avoids the overflow

What actually happened at step 618:

1. The model got confident $\rightarrow$ logits grew past ± 89 in float32
2. $\exp(\text{logit}) \rightarrow \text{inf} \rightarrow \text{inf} / \text{inf} \rightarrow \text{nan}$ in the softmax
3. NaN loss $\rightarrow$ NaN gradients $\rightarrow$ NaN weights, **everywhere, permanently**

This is why PyTorch's `CrossEntropyLoss` takes **raw logits**, not probabilities: it fuses log-softmax into the loss and applies *exactly* the subtract-the-max trick you just derived. The fix for a billion-dollar training run is one line of first-year algebra.

Precision in deep learning


The four formats that run the AI world

Format	Bits	Exp	Mant	Max value	Machine ϵ	~Digits
float64	64	11	52	1.8×10^{308}	2.2×10^{-16}	15.9
float32	32	8	23	3.4×10^{38}	1.2×10^{-7}	7.2
float16	16	5	10	65 504	9.8×10^{-4}	3.3
bfloat16	16	8	7	3.4×10^{38}	7.8×10^{-3}	2.4


Halving storage cuts tensor memory and bandwidth; suitable accelerators can also execute 16-bit matrix operations much faster than float32. The exact speedup depends on hardware and workload. But look at float16's max value...

Range vs precision · the 16-bit dilemma

bit budgets — bfloat16 = float32 with the mantissa chopped



why deep learning picked bfloat16



Why many accelerators use bfloat16

float16 spent its bits on precision and starved the exponent — $\text{exp}(12)$ already overflows its 65,504 ceiling. **bfloat16 keeps float32's 8 exponent bits** and pays with mantissa:

	float16	bfloat16
$\text{exp}(12) \approx 162,755$	inf — training dies	$1.63 \times 10^5 \checkmark$
copy a float32's range?	no — overflow risk everywhere	yes — same exponent, just chop
digits of precision	3.3	2.4

Mixed precision keeps critical work in float32

Low-precision training is not “make every number bfloat16”:

- multiply large tensors in float16 or bfloat16 for speed,
- accumulate dot products and keep selected state in float32,
- use loss scaling when float16 gradients risk underflow.

bfloat16 keeps float32's range; float16 keeps more fraction bits. Mixed precision chooses where each trade-off is acceptable.

Predict all three, then commit.

A. `np.exp(np.float16(12.0))`
returns...?

C. From the start of class: is `0.5 + 0.25 == 0.75` True or False?

B. `np.float32(16_777_216) + np.float32(1.0)` returns...?

D. (bonus) why is `nan == nan` False?

- A.** $\text{inf} - e^{12} \approx 162,755 > 65,504$, float16's ceiling. (bfloat16 shrugs: $\approx 1.6 \times 10^5$.)
- B.** 16777216.0 — unchanged. At 2^{24} the float32 gap is 2, so +1 rounds back down.
- C. True!** $0.5 = 2^{-1}$, $0.25 = 2^{-2}$, $0.75 = 2^{-1} + 2^{-2}$ — all *exactly* representable in binary. Floats aren't sloppy; they are **exact binary fractions**. Trouble only starts when your decimal (0.1) has no finite binary form.
- D.** NaN means “no defensible value” — two undefined results can't be declared equal. $x \neq x$ is the classic NaN test.

Practice problems

Try on paper; verify in the Tutorial 1 notebook.

1. Encode -0.75 as float32 (sign, exponent field, first mantissa bits).
2. What is the largest odd integer float32 represents exactly? (Hint: where does the gap reach 2?)
3. Compute $\text{logsumexp}([-1000, -1001])$ by hand. What would the naive computation return?
4. Explain why $(0.1 + 0.2) + 0.3 \neq 0.1 + (0.2 + 0.3)$ in float64 — what algebraic law do floats break?
5. Predict the sign and rough size of the error when the one-pass variance formula meets $[1e6, 1e6 + 1, 1e6 + 2]$ in float32.
6. Your float16 training run dies at $\exp(12)$. A colleague proposes bfloat16. What exactly changes, and what do you give up?



Kahan summation

★ OPTIONAL

Add `np.float32(0.1)` ten million times. True answer: 10^6 . A naive loop returns...

1087937.0 — off by 8.8%

Once the sum reaches 10^6 , the grid gap is 0.0625: each $+0.1$ rounds badly, ten million times, *in the same direction*.

```
s, c = 0.0, 0.0
for x in xs:
    y = x - c           # re-inject the error we still owe
    t = s + y           # big + small: low digits of y get dropped...
    c = (t - s) - y     # ...(t - s) - y recovers exactly what was dropped
    s = t
```

What compensated summation guarantees

Kahan's variable c carries low-order bits that the running sum lost. This greatly reduces first-order accumulation error, especially when many same-sign terms are added.

The error still depends on the conditioning of the sum and on higher-order terms; it is not literally independent of n . NumPy commonly reduces error with a different method, **pairwise summation**.



Subnormals: the fine print near zero

★ OPTIONAL

The smallest *normal* float32 is $2^{-126} \approx 1.18 \times 10^{-38}$. Without further tricks there'd be a **moat around zero**, and $x - y == 0$ could be true for $x \neq y(!)$.

Subnormals: when the exponent field is all zeros, drop the implicit leading 1 and let the mantissa alone carry the value — filling the moat with evenly spaced numbers down to $2^{-149} \approx 1.4 \times 10^{-45}$ (“gradual underflow”).

The catch: some hardware handles subnormals in microcode — **orders of magnitude slower**. GPUs and DL runtimes often enable *flush-to-zero* modes that trade this correctness for speed. If a kernel inexplicably slows down as values shrink toward zero: suspect subnormals.

Lecture 2 — summary

- **Fixed formats fail ML**: one program spans 40+ orders of magnitude → the point must float.
- **IEEE-754 float32** = sign + biased exponent (127) + mantissa with a free leading 1; 6.25 encodes exactly, 0.1 *cannot* — hence $0.1 + 0.2 \neq 0.3$.
- **The float line is unevenly spaced**: gaps double at each power of 2; $\varepsilon_{32} \approx 10^{-7}$; at 2^{24} , $x + 1 == x$.
- **Special values**: overflow → `inf`, indeterminate → `nan`; NaN propagates and `nan != nan`.
- **Catastrophic cancellation**: subtracting near-equals shreds digits — rewrite the formula.
- **The ML fixes**: log-space probabilities, log-sum-exp, subtract-the-max softmax.
- **bf16** = float32's range at half the bits: DL trades precision for range.

Before L3 and Tutorial 1

- **Reading:** Goldberg (1991), §1–2; Solomon, *Numerical Algorithms*, Ch. 1–2.
- **Worksheet:** decode IEEE-754 values by hand.
- **Notebook:** reproduce decimal error, cancellation, and softmax overflow in NumPy, then implement the stable versions.

Computers don't do real numbers — they do ~4 billion of them, unevenly spaced.

Next: **vectors** — how direction and displacement explain $\text{king} - \text{man} + \text{woman} \approx \text{queen}$.