

Jacobian, Hessian & Multivariate Taylor

When derivatives become matrices


Prof. Nipun Batra

Mathematical Foundations for AI · IIT Gandhinagar

The hook: a derivative with a million inputs

The function AI actually trains

A neural network with its loss attached is **one function**: parameters in, a single number out.



For real models $d \approx 10^6$ to 10^{12} . Training (L17) needs this function's **derivative**.

But what single object is “the derivative” of $\mathbb{R}^{10^6} \rightarrow \mathbb{R}$? And between the layers sit maps $\mathbb{R}^n \rightarrow \mathbb{R}^m$ — vectors in, **vectors out**. L7's one number cannot be the answer.

Today's promise — and today's safety rail

The derivative grows up in two steps:

- **vector in, vector out** → the derivative is a matrix: the **Jacobian**;
- **curvature** of a surface → also a matrix: the **Hessian** — and its eigenvalue signs (L5!) say whether you're in a bowl, on a saddle, or in a trough.

Fair warning: this is the **heaviest-notation lecture of the course** — so we run it with a safety valve. Every idea happens on $\mathbb{R}^2 \rightarrow \mathbb{R}^2$, small enough to check by hand. The general case is always the **same picture, more rows**.

By the end, “Jacobian” and “Hessian” will be **pictures**, not notation: a parallelogram, and a bowl.

One table to navigate the whole lecture. Two rows you own; two we fill today:

function	derivative	shape	what it is
$f : \mathbb{R} \rightarrow \mathbb{R}$	$f'(x)$	a number	slope of the local line (L7)
$f : \mathbb{R}^n \rightarrow \mathbb{R}$	$\nabla f(x)$	vector, n	all partials; points straight uphill (L8)
$f : \mathbb{R}^n \rightarrow \mathbb{R}^m$	?	?	today
$f : \mathbb{R}^n \rightarrow \mathbb{R}$, order 2	?	?	today

Same idea every row: **the derivative is the best local linear story** — only its container grows.

Learning outcomes


By the end of this lecture you will be able to:

1. Compute a 2×2 **Jacobian**, read it as a local linear map, and read $|\det J|$ as area scaling.
2. Assemble a **Hessian** and explain entry (i, j) as “how slope i responds to nudge j ”.
3. Classify $x^\top A x$ as bowl / saddle / trough from **eigenvalue signs**, with 2×2 hand tests.
4. Derive ★ $\nabla(x^\top A x) = (A + A^\top)x$.
5. Write the **second-order Taylor expansion** — the formula every optimizer lives on.
6. Keep shapes straight: gradient-as-column, J is $m \times n$, chain rule = matrix product.

Vector in, vector out: the Jacobian

A function that moves points

$\Phi : \mathbb{R}^2 \rightarrow \mathbb{R}^2$ eats a point, returns a point. Feed it a whole grid and watch space bend:



$\Phi(x, y) = (x + 0.6 \sin 1.6y, y + 0.6 \sin 1.6x)$ — and every network layer is one of these (L4)

Globally: curvy. But the little orange square just got **leaned and stretched**. Hold that thought.

The one case we already own: a linear map

For a **linear** map (L4: “a matrix is a verb”) the nudge story is exact, everywhere, with no zooming:

$$F(x) = Ax = \begin{pmatrix} 2 & 1 \\ 3 & -2 \end{pmatrix} x \Rightarrow F(a + \delta) = F(a) + A\delta$$

Nudge the input by δ — the output moves by $A\delta$. **The matrix A is the derivative**, at every point.

Target for today: for a curvy Φ , find the matrix that plays the role of A near one point.

Numbers: nudge the polar map, watch both outputs

Take the most useful curvy $\mathbb{R}^2 \rightarrow \mathbb{R}^2$ map there is — polar coordinates — at $(r, \theta) = (2, \frac{\pi}{3})$:

$$\Phi(r, \theta) = \left(\underbrace{r \cos \theta}_x, \underbrace{r \sin \theta}_y \right), \quad \Phi\left(2, \frac{\pi}{3}\right) = (1.000, 1.732)$$

Nudge each input by 0.01, one at a time (L7's nudge machine, run twice per output):

nudge	Δx , per unit nudge	Δy , per unit nudge
$r : 2 \rightarrow 2.01$	$+0.00500/0.01 = 0.500$	$+0.00866/0.01 = 0.866$
$\theta : \frac{\pi}{3} \rightarrow \frac{\pi}{3} + 0.01$	$-0.01737/0.01 = -1.737$	$+0.00991/0.01 = 0.991$

Two inputs $\times$ two outputs = **four** sensitivities. One number was never going to be enough — but four numbers have an obvious home.

Pack the four sensitivities: the Jacobian

House rule: **one row per output, one column per input**. Each entry is a partial derivative (L8):


$$J = \begin{pmatrix} \frac{\partial x}{\partial r} & \frac{\partial x}{\partial \theta} \\ \frac{\partial y}{\partial r} & \frac{\partial y}{\partial \theta} \end{pmatrix} = \begin{pmatrix} \cos \theta & -r \sin \theta \\ \sin \theta & r \cos \theta \end{pmatrix} \xrightarrow{(r, \theta) = (2, \pi/3)} \begin{pmatrix} 0.500 & -1.732 \\ 0.866 & 1.000 \end{pmatrix}$$

Exactly the four measured ratios (the -1.737 and 0.991 were drifting here as the nudge shrank).

- Row 1 = $\nabla x^\top$: how output x feels each input — a gradient lying on its side.
- Sanity check on the linear map: every partial of Ax is a constant, and $J = A$. ✓

The picture: little squares become parallelograms

True image of a tiny square (ink) vs the parallelogram J makes of it (orange) — the mismatch **dies faster than ε** :




$$\Phi(\mathbf{a} + \delta) \approx \Phi(\mathbf{a}) + J(\mathbf{a}) \delta$$

The Jacobian is the linear map that impersonates Φ near $\mathbf{a}$.

$\det J$: the area dial

L4: a determinant is **the area-scaling factor of a linear map** — so $|\det J|$ is Φ 's local area dial:



For polar: $\det J = r \cos^2 \theta + r \sin^2 \theta = r$ — the orange cell ($r = 1.5$) gets $\sim 3 \times$ the teal cell's area ($r = 0.5$). At $r = 0$: $\det J = 0$, the whole θ -axis is **crushed to a point** (L4's null space, gone local).

Code: PyTorch measures the same matrix

```
import torch
from torch.autograd.functional import jacobian

def polar(p):
    r, th = p
    return torch.stack([r * torch.cos(th), r * torch.sin(th)])

J = jacobian(polar, torch.tensor([2.0, torch.pi / 3]))
# tensor([[ 0.5000, -1.7321],
#         [ 0.8660,  1.0000]])
torch.det(J)          # tensor(2.0000) = r * r
```

Our four hand-measured ratios, to four decimals — and the area dial reads $r = 2$, as derived.

Planted for L12: when probability densities change variables, they must pay exactly this $|\det J|$ area toll. Same dial, new customer.

The general case: same picture, more rows

For $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$ — nothing new, just a bigger crate (m outputs stacked, n inputs across):

$$J(\mathbf{x}) = \begin{pmatrix} \frac{\partial f_1}{\partial x_1} & \cdots & \frac{\partial f_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_m}{\partial x_1} & \cdots & \frac{\partial f_m}{\partial x_n} \end{pmatrix} \in \mathbb{R}^{m \times n}, \quad J_{ij} = \text{how output } i \text{ feels input } j$$

The zoo collapses out of this one object:

- $m = 1$: a single row — the gradient lying on its side, $J = \nabla f^\top$ (L8).
- $m = n = 1$: a 1×1 matrix — L7's lone number $f'(x)$.

There is only **one derivative idea. The Jacobian is its general form; slope and gradient are its small print.**

Chain rule: composition becomes multiplication

L7's chain rule multiplied slopes. Jacobians **are** the local linear maps — and composing linear maps is matrix multiplication (L4):

$$J_{g \circ f}(\mathbf{a}) = \underbrace{J_g(f(\mathbf{a}))}_{k \times m} \cdot \underbrace{J_f(\mathbf{a})}_{m \times n} \in \mathbb{R}^{k \times n}$$

Square through $f \rightarrow$ parallelogram; that parallelogram through $g \rightarrow$ another parallelogram. Multiply the matrices, compose the leans.

And the shapes **must click**: $(k \times m)(m \times n) = (k \times n)$ — your free error-detector for every derivation.

Planted for L10–L11: a network is $f_L \circ \dots \circ f_1$, so its derivative is a **product of Jacobians**. Backprop will be nothing but a clever order for multiplying them out.

function	derivative	shape	what it is
$f : \mathbb{R} \rightarrow \mathbb{R}$	$f'(x)$	a number	slope of the local line (L7)
$f : \mathbb{R}^n \rightarrow \mathbb{R}$	$\nabla f(x)$	vector, n	all partials; points straight uphill (L8)
$f : \mathbb{R}^n \rightarrow \mathbb{R}^m$	$J(x)$	matrix, $m \times n$	the local linear map (today)
$f : \mathbb{R}^n \rightarrow \mathbb{R}$, order 2	?	?	today

One row left. It needs a different question: not “vector out?” but “**what happened to curvature?**”

$f(x, y, z) = (x^2 + y^2, y - z)$ maps $\mathbb{R}^3 \rightarrow \mathbb{R}^2$. What is the shape of its Jacobian — and is $\det J$ available to talk about areas?

A. 3×2 , and yes

B. 2×3 , and no

C. 2×3 , and yes

D. 3×3 , and yes

Correct: B — J is 2×3 , and $\det J$ does not exist

Why: One row per **output** (2), one column per **input** (3): $J = \begin{pmatrix} 2x & 2y & 0 \\ 0 & 1 & -1 \end{pmatrix}$, a 2×3 matrix. Determinants belong to **square** matrices only — a map $\mathbb{R}^3 \rightarrow \mathbb{R}^2$ flattens 3-D volumes onto a plane, and “how much did the volume scale?” stops being a fair question. The area differential needs $n = m$.

Curvature becomes a matrix: the Hessian

In one dimension, curvature was one number

L7's second act: at a flat point, the sign of f'' told the whole story — smile $\rightarrow$ minimum, frown $\rightarrow$ maximum.

Now stand on a 2-D surface $f : \mathbb{R}^2 \rightarrow \mathbb{R}$ and walk in different directions:

- along one street the ground may **curve up** like a smile,
- along the cross-street it may curve up harder, or not at all — **or curve down**.

One number per direction, infinitely many directions. Sounds hopeless — but so did “one slope per direction” before L8 packed every directional slope into **one vector** ∇f .

Same trick again, one storey up.

Numbers: the gradient is itself a map we know

Take $f(x, y) = x^2 + xy + y^2$ (keep it in view — it returns all lecture). Its gradient:

$$\nabla f(x, y) = \begin{pmatrix} 2x + y \\ x + 2y \end{pmatrix}$$

Look at what that **is**: vector in, vector out — $\nabla f : \mathbb{R}^2 \rightarrow \mathbb{R}^2$. We built the perfect tool for those **ten minutes ago**.

Take the Jacobian **of the gradient**:

$$J_{\nabla f} = \begin{pmatrix} \frac{\partial}{\partial x}(2x + y) & \frac{\partial}{\partial y}(2x + y) \\ \frac{\partial}{\partial x}(x + 2y) & \frac{\partial}{\partial y}(x + 2y) \end{pmatrix} = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$$

That matrix has a name — and for this f , it is **L5's favorite matrix**. Not a coincidence we'll waste.

The Hessian, defined

$$H(\mathbf{x}) = J_{\nabla f}(\mathbf{x}), \quad H_{ij} = \frac{\partial^2 f}{\partial x_i \partial x_j} = \text{how slope } i \text{ responds to a nudge of } x_j$$

- Diagonal H_{ii} : L7's old curvature along axis i — walk east, does the eastward slope grow?
- Off-diagonal H_{ij} : the **twist** — walk north, does the **eastward** slope change?

And a gift: for any twice-continuously-differentiable f , mixed partials agree — $H_{ij} = H_{ji}$ (Schwarz's theorem, stated with a straight face like L5's spectral theorem).

The Hessian is **symmetric — so the spectral theorem applies: real eigenvalues, perpendicular eigenvectors. L5's whole toolkit just walked in the door.**

Test drive: the curvature you feel along a street

Walking from a in unit direction u , the ground under you is the 1-D function $g(t) = f(a + tu)$, and (chain rule, twice):

$$g''(0) = u^\top H u$$

For our f with $H = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$, feel three streets:

direction u	$u^\top H u$	the street feels like
$(1, 0)$	2	a smile
$(1, 1)/\sqrt{2}$	$(2 + 1 + 1 + 2)/2 = 3$	the steepest smile
$(1, -1)/\sqrt{2}$	$(2 - 1 - 1 + 2)/2 = 1$	the gentlest smile

3 and 1 are H 's **eigenvalues**, felt along its **eigenvectors** (L5: $\lambda = 3, 1$ at $(1, 1), (1, -1)$). The extreme curvatures of a surface are an eigen-problem.

Same picture, more rows: a 3-variable Hessian

$f(x, y, z) = x^2 + y^2 + xyz$. First the gradient, then differentiate **it**, coordinate by coordinate:

$$\nabla f = \begin{pmatrix} 2x + yz \\ 2y + xz \\ xy \end{pmatrix} \Rightarrow H = \begin{pmatrix} 2 & z & y \\ z & 2 & x \\ y & x & 0 \end{pmatrix} \xRightarrow{(2,1,3)} \begin{pmatrix} 2 & 3 & 1 \\ 3 & 2 & 2 \\ 1 & 2 & 0 \end{pmatrix}$$

Symmetric — check. Entries depend on where you stand — curvature is **local**, computed per point.

Read $H_{12} = z$ out loud: “how the x -slope responds to nudging y .” At $(2, 1, 3)$: nudge y by 0.01 and the x -slope $2x + yz$ climbs by 0.03. Every entry of every Hessian is a sentence like this.

Code: the Hessian in one call

```
from torch.autograd.functional import hessian

f = lambda v: v[0]**2 + v[0]*v[1] + v[1]**2
hessian(f, torch.tensor([1.0, 1.0]))
# tensor([[2., 1.],
#         [1., 2.]])

g = lambda v: v[0]**2 + v[1]**2 + v[0]*v[1]*v[2]
hessian(g, torch.tensor([2.0, 1.0, 3.0]))
# tensor([[2., 3., 1.], [3., 2., 2.], [1., 2., 0.]])
```

Under the hood it does what we did: differentiate the gradient — the Jacobian of ∇f .

The shape gallery: quadratic forms

The pure-curvature family



Which functions have the **same** Hessian everywhere — curvature standing still so we can stare at it? Quadratics. For a symmetric A , multiplied out once by hand on 2×2 :

$$q(\mathbf{x}) = \frac{1}{2} \mathbf{x}^\top A \mathbf{x} = \frac{1}{2} (x_1, x_2) \begin{pmatrix} a & b \\ b & c \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \frac{1}{2} (ax_1^2 + 2bx_1x_2 + cx_2^2) \quad \text{— a number}$$



For $A = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$: $q = x^2 + xy + y^2$ — **our running f** . Claim (★ proof in the next section): $\nabla q = Ax$ and $H = A$, at every point.

A quadratic form is a Hessian standing still. Classify its shapes, and you've classified **all local curvature.**



The gallery: bowls, saddles, troughs



$\lambda = 3, 1$ – both > 0



$\lambda = 2, -2$ – mixed signs



$\lambda = 2, 0$ – a zero

surfaces and their contour maps, computed live from $q(x) = \frac{1}{2}x^\top Ax$; eigenvalues by the in-slide solver – ellipses, hyperbolas, parallel lines

Eigenvalues read the shape — L5, cached in

A is symmetric, so L5 fires: $A = Q\Lambda Q^\top$ — “a perpendicular grid of dials”. In eigen-coordinates $t = Q^\top x$ the cross-term dies; the dials are **curvatures**:

$$q = \frac{1}{2}(\lambda_1 t_1^2 + \lambda_2 t_2^2 + \cdots + \lambda_n t_n^2) \quad \text{— one independent parabola per eigen-direction}$$

eigenvalue signs	shape	at a critical point this means
all $\lambda_i > 0$	bowl	strict local minimum — every street smiles
all $\lambda_i \geq 0$, some = 0	trough	a flat valley floor along the $\lambda = 0$ directions
mixed signs	saddle	up some streets, down others — not a min!
all $\lambda_i < 0$	dome	local maximum

Planted for L13: the bowl’s ellipse contours (axes = eigenvectors) are exactly a Gaussian’s, with $A = \Sigma^{-1}$.

The official words: definiteness

The vocabulary every ML paper uses for “which shape is it”:

A is...	definition: for all $x \neq 0$	equivalently
positive definite (PD)	$x^\top A x > 0$	all $\lambda_i > 0$
positive semi-definite (PSD)	$x^\top A x \geq 0$	all $\lambda_i \geq 0$
negative definite (ND)	$x^\top A x < 0$	all $\lambda_i < 0$
indefinite	both signs occur	mixed-sign λ_i

Why the two columns agree: in eigen-coordinates $x^\top A x = \sum_i \lambda_i t_i^2$ — a sum of parabolas can only be always-positive if every dial λ_i is.

Positive definite = the surface is a bowl. That one sentence decodes half the theorems you will ever meet about optimization and covariance.

Hand tests for 2×2 — ten-second shape reading

L5's sanity checks were $\text{tr } A = \sum \lambda_i$ and $\det A = \prod \lambda_i$. For symmetric 2×2 they **decide the shape**:

test	verdict
$\det A < 0$	λ 's have opposite signs — saddle , done (a negative det always betrays a flip: L4, L5)
$\det A > 0$	same signs — read $\text{tr } A$: positive $\rightarrow$ bowl , negative $\rightarrow$ dome
$\det A = 0$	some $\lambda = 0$ — semi-definite : trough (or flat)

For $n \times n$, the computer's test is eigenvalues; the classical hand test is **Sylvester's criterion** (all leading principal minors $> 0 \Leftrightarrow$ PD — stated for your shelf, not proved).

Positive **entries** do not make a matrix positive **definite**: $\begin{pmatrix} 1 & 2 \\ 2 & 1 \end{pmatrix}$ is all-positive yet $\det = -3 < 0$ — a saddle. Definiteness lives in the eigenvalues, not the ink.

Numbers: classify three matrices

A	tr	det	verdict by hand	λ (live)
$\begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$	4	3	det > 0, tr > 0 — bowl (PD)	3, 1
$\begin{pmatrix} 1 & 2 \\ 2 & 1 \end{pmatrix}$	2	−3	det < 0 — saddle (indefinite)	3, −1
$\begin{pmatrix} 1 & 1 \\ 1 & 1 \end{pmatrix}$	2	0	det = 0, tr > 0 — trough (PSD)	2, 0

The λ column is computed by the in-slide eigensolver at compile time — trace and det agree with $\sum \lambda_i$ and $\prod \lambda_i$ on every row (L5's exam armor, still working).

Code: three lines settle any matrix

```
import numpy as np

for A in ([[2, 1], [1, 2]], [[1, 2], [2, 1]], [[1, 1], [1, 1]]):
    lam = np.linalg.eigvalsh(A)      # symmetric -> real, sorted
    print(lam)
# [1. 3.]      all > 0          -> bowl      (PD)
# [-1. 3.]     mixed signs     -> saddle    (indefinite)
# [0. 2.]      a zero, rest +  -> trough   (PSD)
```

eigvalsh — the h means **Hermitian/symmetric**: it promises real eigenvalues because L5's spectral theorem promises them first.

At a critical point ($\nabla f = 0$), the Hessian is $H = \begin{pmatrix} 4 & -2 \\ -2 & 1 \end{pmatrix}$. The surface there is a...

A. bowl — strict local minimum

B. saddle

C. trough — a flat direction

D. dome — local maximum

Correct: C — trough — a flat valley direction

Why: $\det H = 4 - 4 = 0$, so a λ is exactly 0; $\text{tr } H = 5 > 0$ pins the other at +5 (indeed $\lambda^2 - 5\lambda + 0 = 0 \Rightarrow \lambda = 5, 0$). Positive curvature along one eigen-street, **perfectly flat** along the other $((1, 2), \text{ where } Hu = 0)$ — a PSD trough, not a strict minimum. Beware option A: $\det = 0$ is precisely where “bowl” claims go to die.

The one derivation: the gradient of a quadratic form

Why this gradient, of all gradients

$x^\top Ax$ is the most-differentiated expression in machine learning:

- least squares (L4's preview): $\|X\theta - y\|^2$ expands into one;
- the Gaussian's exponent (L13) **is** one; the L2 regularizer $\theta^\top \theta$ is one with $A = I$;
- every second-order Taylor model (ten minutes from now) ends in one.

So we earn its gradient honestly — today's , the difficulty valve fully open:

Plan: brute-force the 2×2 case symbol by symbol (no matrix calculus allowed), **then** watch the general answer repack itself. One derivation, everything else today is plausibility.

★ Step 1: multiply the 2×2 out, no shortcuts

D DERIVATION

Let $A = \begin{pmatrix} a & b \\ c & d \end{pmatrix}$ (deliberately **not** symmetric) and $x = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$:

$$x^\top A x = (x_1, x_2) \begin{pmatrix} ax_1 + bx_2 \\ cx_1 + dx_2 \end{pmatrix}$$

$$= ax_1^2 + bx_1x_2 + cx_1x_2 + dx_2^2 = ax_1^2 + (b + c)x_1x_2 + dx_2^2$$

A plain two-variable polynomial — L8 can differentiate this in its sleep. Note who survived: b and c only ever appear as the **sum** $b + c$.

★ Step 2: two honest partial derivatives

$$\frac{\partial}{\partial x_1} [ax_1^2 + (b+c)x_1x_2 + dx_2^2] = 2ax_1 + (b+c)x_2$$

$$\frac{\partial}{\partial x_2} [ax_1^2 + (b+c)x_1x_2 + dx_2^2] = (b+c)x_1 + 2dx_2$$

Stack them into the gradient (our column convention):

$$\nabla(\mathbf{x}^\top A \mathbf{x}) = \begin{pmatrix} 2ax_1 + (b+c)x_2 \\ (b+c)x_1 + 2dx_2 \end{pmatrix}$$

★ Step 3: recognize the matrix hiding in it

D DERIVATION

$$\begin{pmatrix} 2ax_1 + (b+c)x_2 \\ (b+c)x_1 + 2dx_2 \end{pmatrix} = \begin{pmatrix} 2a & b+c \\ b+c & 2d \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = \left[\begin{pmatrix} a & b \\ c & d \end{pmatrix} + \begin{pmatrix} a & c \\ b & d \end{pmatrix} \right] \mathbf{x}$$

$$\nabla(\mathbf{x}^\top A \mathbf{x}) = (A + A^\top) \mathbf{x} \quad \Rightarrow \quad \text{symmetric } A : \quad \nabla = 2A \mathbf{x}$$

Same argument in $\mathbb{R}^n$, one line: $\mathbf{x}^\top A \mathbf{x} = \sum_{i,j} A_{ij} x_i x_j$, and $\partial/\partial x_k$ harvests row k ($\sum_j A_{kj} x_j$) plus column k ($\sum_i A_{ik} x_i$) — that is $(A \mathbf{x})_k + (A^\top \mathbf{x})_k$. Same picture, more rows.

Corollary (one more derivative): $H = J_\nabla = A + A^\top$ — for symmetric A and $q = \frac{1}{2} \mathbf{x}^\top A \mathbf{x}$: $\nabla q = A \mathbf{x}$, $H = A$. The gallery's caption is now a theorem.

The rule, cross-examined

The slide checks itself. Chalkdust's autodiff evaluates $\nabla(x^\top Ax)$ for $A = \begin{pmatrix} 2 & 1 \\ 1 & 2 \end{pmatrix}$ at $(1, 1)$ — reverse mode, no symbols:

$$\text{autodiff: } \nabla = \begin{pmatrix} 6.0000 \\ 6.0000 \end{pmatrix} \quad \text{our rule: } 2A \begin{pmatrix} 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 6.0000 \\ 6.0000 \end{pmatrix} \quad \text{— identical } \checkmark$$

Two rules you now own for free, by choosing A :

expression	gradient	why
$\ x\ ^2 = x^\top Ix$	$2x$	$A = I$ is symmetric
$b^\top x$ (linear)	b	the warm-up act: no square, constant slope

These two + today's  generate most of the Matrix Cookbook's first page. The rest is looked up, not re-derived — more on that in the bookkeeping section.

Taylor, one degree at a time

L7's ladder, one more time

L7 replaced a function near x_0 by polynomials, one degree at a time:

$$f(x_0 + h) \approx f(x_0) + f'(x_0) h + \frac{1}{2} f''(x_0) h^2$$

Every ingredient just grew up. Upgrade each term with today's zoo:

1-D ingredient	n -D upgrade	the term becomes
step h	step vector δ	
$f'(x_0) h$	∇f (L8)	$\nabla f(x)^\top \delta$ — a dot product
$\frac{1}{2} f''(x_0) h^2$	H (today)	$\frac{1}{2} \delta^\top H(x) \delta$ — a quadratic form

Slope became a vector, curvature became a matrix — and h^2 became “sandwich H between two δ 's”.

The formula every optimizer lives on

$$f(\mathbf{x} + \boldsymbol{\delta}) \approx \underbrace{f(\mathbf{x})}_{\text{height}} + \underbrace{\nabla f(\mathbf{x})^\top \boldsymbol{\delta}}_{\text{tilt: the local plane}} + \underbrace{\frac{1}{2} \boldsymbol{\delta}^\top H(\mathbf{x}) \boldsymbol{\delta}}_{\text{bend: the local bowl or saddle}}$$

Read it as a story: stand at $\mathbf{x}$; the surface near you is **a height, plus a tilt, plus a pure-curvature shape from the gallery** — and you can now classify that shape by eigenvalue signs.

L16–L18 run on this one line: gradient descent trusts the tilt, Newton trusts the bend. Today it just has to become yours.

Numbers: predict a function you've never met

$f(x, y) = x^2y$ at $\mathbf{a} = (1, 1)$, step $\delta = (0.1, 0.1)$. Harvest the three ingredients:

$$f(\mathbf{a}) = 1, \quad \nabla f = \begin{pmatrix} 2xy \\ x^2 \end{pmatrix} = \begin{pmatrix} 2 \\ 1 \end{pmatrix}, \quad H = \begin{pmatrix} 2y & 2x \\ 2x & 0 \end{pmatrix} = \begin{pmatrix} 2 & 2 \\ 2 & 0 \end{pmatrix}$$

Assemble, term by term:


$$f(\mathbf{a} + \delta) \approx 1 + \underbrace{(2, 1) \cdot (0.1, 0.1)}_{0.3} + \underbrace{\frac{1}{2}(0.1, 0.1) \begin{pmatrix} 2 & 2 \\ 2 & 0 \end{pmatrix} \begin{pmatrix} 0.1 \\ 0.1 \end{pmatrix}}_{\frac{1}{2}(0.06)=0.03} = 1.33$$

Truth: $f(1.1, 1.1) = 1.1^2 \times 1.1 = 1.331$. Off by 0.001 — exactly the dropped third-order term.

Note the Hessian is **indefinite** here ($\det = -4 < 0$: saddle curvature) — Taylor doesn't care; it reports the local shape, whatever it is.

The three stories, along one street

Walk from $a = (1, 1)$ along $\delta = (t, t)$: the true f becomes $g(t) = (1 + t)^3$, and the Taylor terms become L7-style line and parabola — all three plotted live:



Near a : the parabola hugs the truth visibly longer than the line — the multivariate story, sliced back down to L7's picture.

Code: the whole expansion in five lines

```
from torch.autograd.functional import jacobian, hessian

f = lambda v: v[0]**2 * v[1]
a = torch.tensor([1.0, 1.0]); d = torch.tensor([0.1, 0.1])

g, H = jacobian(f, a), hessian(f, a)
pred2 = f(a) + g @ d + 0.5 * d @ H @ d
# pred2 = 1.3300      true f(a + d) = 1.3310      gap = 3rd order
```

Five lines you will reuse in T5, and conceptually in every optimizer of Module 4.

A promissory note: where this formula gets cashed

lecture	what it does with today's expansion
L16	convexity : H PSD everywhere — the surface is one global bowl, no bad valleys
L17	gradient descent : trust the tilt, step downhill; the learning rate is literally “how far you trust the linear term” — and H 's eigenvalues decide its fate (L5's warning)
L18	Newton's method : trust the bowl — minimize the quadratic model, jump $\delta^* = -H^{-1}\nabla f$

Every optimizer is a policy about **which Taylor term to trust, and how far.**

Bookkeeping & the take-home

Our contract: the gradient is a column

Matrix calculus has a dirty secret: two bookkeeping schools. **Numerator layout** writes the derivative of a scalar as a row; **denominator layout** as a column. Both are fine; mixing them mid-derivation is how homework dies.

This course signs one contract and sticks to it (MML's):

object	shape	note
$x, \delta, \nabla f$	column, n	gradient-as-column — same shape as x , so $x - \eta \nabla f$ type-checks (L17)
J of $f : \mathbb{R}^n \rightarrow \mathbb{R}^m$	$m \times n$	rows = outputs; for $m = 1$, $J = \nabla f^\top$ (a row)
H	$n \times n$	symmetric; $H = J_{\nabla f}$

When you open any other book or blog, **find its convention first** — a transposed-looking formula is usually the other school, not an error.

Shape-checking is a proof technique

The chain rule **forces** every shape, so let the shapes catch your bugs:

$$\underbrace{J_{g \circ f}}_{k \times n} = \underbrace{J_g}_{k \times m} \underbrace{J_f}_{m \times n}, \quad \underbrace{\nabla f(\mathbf{x})^\top \boldsymbol{\delta}}_{(1 \times n)(n \times 1) = \text{scalar}}, \quad \underbrace{\boldsymbol{\delta}^\top}_{1 \times n} \underbrace{H}_{n \times n} \underbrace{\boldsymbol{\delta}}_{n \times 1} = \text{scalar}$$

If a derivation produces a shape mismatch, it is wrong — no further reading required. That five-second check replaces half of matrix-calculus memorization.

References for life: the Matrix Cookbook (look identities up — nobody re-derives them), Parr & Howard, *The Matrix Calculus You Need for Deep Learning* (the gentle full story), MML §5.4–5.7 (our text, same conventions).

The hook, resolved

A neural network is a chain $\mathcal{L} = f_L \circ \dots \circ f_1(\theta)$. What is its derivative? Now we can **name every piece**:

1. each layer $f_i : \mathbb{R}^n \rightarrow \mathbb{R}^m$ has a **Jacobian** — a lean-and-stretch parallelogram map;
2. the chain rule multiplies them: $J = J_L \dots J_1$;
3. the last output is the scalar loss, so the whole product collapses to **one row** — the gradient $\nabla \mathcal{L}^\top$, L8's uphill vector, for a million inputs;
4. the loss's curvature is a **Hessian** — but at $d = 10^6$ it has 10^{12} entries (≈ 4 TB in float32): real optimizers never store it, they respect it (L17–L18).

Promise kept: the Jacobian is a parallelogram, the Hessian is a bowl — and L10–L11 build the machine that multiplies those Jacobians for you.

function	derivative	shape	what it is
$f : \mathbb{R} \rightarrow \mathbb{R}$	$f'(x)$	a number	slope of the local line (L7)
$f : \mathbb{R}^n \rightarrow \mathbb{R}$	$\nabla f(x)$	vector, n	all partials; points straight uphill (L8)
$f : \mathbb{R}^n \rightarrow \mathbb{R}^m$	$J(x)$	matrix, $m \times n$	the local linear map (today)
$f : \mathbb{R}^n \rightarrow \mathbb{R}$, order 2	$H(x)$	matrix, $n \times n$, symmetric	curvature, in every direction (today)

One idea, four containers: **the derivative is the best local linear (then quadratic) story** — the container just matches the function's shape.

Lecture 9 — summary

- **Jacobian**: $m \times n$ crate of partials; the local linear map — $\Phi(\mathbf{a} + \delta) \approx \Phi(\mathbf{a}) + J\delta$; little squares $\rightarrow$ parallelograms; $|\det J|$ = the area dial (polar: $\det J = r$).
- **Chain rule** = matrix multiplication of Jacobians; shapes must click $((k \times m)(m \times n))$.
- **Hessian** = Jacobian of the gradient; symmetric; entry (i, j) = “how slope i feels nudge j ”; $\mathbf{u}^\top H \mathbf{u}$ = curvature along $\mathbf{u}$, extremes at eigenvectors.
- **Quadratic forms** $\frac{1}{2} \mathbf{x}^\top A \mathbf{x}$: bowls (all $\lambda > 0$), saddles (mixed), troughs (a zero λ); 2×2 hand test via det and trace; ★ $\nabla(\mathbf{x}^\top A \mathbf{x}) = (A + A^\top) \mathbf{x}$.
- **Taylor, order 2**: $f + \nabla f^\top \delta + \frac{1}{2} \delta^\top H \delta$ — height, tilt, bend; the optimizer’s home.

Read before L10: MML §5.4–5.7; skim explained.ai/matrix-calculus keep the Matrix Cookbook bookmarked as a reference, not bedtime reading. **T5** (after L11) drills Jacobians and the full backward pass.

Practice problems

Try on paper; verify with `torch.autograd.functional` as in today's code slides.

1. $F(x, y) = (x^2y, x + y^3)$: compute J at $(1, 1)$ and $\det J$. Is F locally inflating or crushing areas there — by what factor?
2. For $\Phi(x, y) = (e^x \cos y, e^x \sin y)$: show $\det J = e^{2x}$. Why can this map never crush area to zero?
3. $f(x, y) = x^3 + y^3 - 3xy$: find both critical points, compute H at each, and classify (one saddle, one bowl).
4. $f(x, y) = xy$ has **no** squared term anywhere. Compute H , its eigenvalues, and the shape at $(0, 0)$.
5. Apply the ★ rule to the non-symmetric $A = \begin{pmatrix} 1 & 4 \\ 0 & 1 \end{pmatrix}$: write $\nabla(x^\top Ax)$, then verify by expanding $x^\top Ax$ by hand.
6. Second-order Taylor of $f(x, y) = e^x \sin y$ at $(0, \frac{\pi}{2})$: show $\nabla f = (1, 0)$, $H = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$, predict $f(0.1, \frac{\pi}{2} + 0.1)$, and compare with the true value 1.0997.



Beyond order two: derivatives become tensors

★ OPTIONAL

Why does everyone stop at the Hessian? Keep differentiating and count indices:

order	object	entries at $d = 10^6$
1	∇f — vector	10^6 (4 MB)
2	H — matrix	10^{12} (4 TB)
3	$\partial^3 f$ — a 3-index tensor	10^{18} (4 EB — exabytes)

Storage explodes by d per order while (L7's lesson) each Taylor term buys less — so ML lives on order 1, consults order 2, and never writes order 3 down.

Escape hatch, teased for L11 : you can compute **Hessian-vector products** Hv without ever materializing H — differentiate $\nabla f^\top v$ once more. Curvature information at gradient prices.

Curvature is a matrix, and its sign is the shape of the bowl.

Next: **Differentiation on a Computer** — the computer takes over differentiation: finite differences pick a fight with floating point (L2 returns), and forward mode wins it.