

Learning Sequence Models

Counts, smoothing, evaluation, generation — the finale

Prof. Nipun Batra

Mathematical Foundations for AI · IIT Gandhinagar

The promise from Lecture 1

Lecture 1 showed this output

Lecture 1 promised a character-level language model, trained from scratch on Shakespeare by L26, and showed its output:

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

CLARENCE:

And so the sight the world and the more,
I shall be my lord the king of hearth.

Today we build that pipeline — count, smooth, score, sample — and run it on Shakespeare before the lecture ends.

Today's pipeline, module by module

1. estimate transition probabilities by counting — maximum likelihood (L14),
2. derive the count formula with Lagrange multipliers (L19),
3. smooth the counts with a conjugate Dirichlet prior (L15),
4. score held-out text with cross-entropy in bits (L24),
5. compute every score in log-space (L2),
6. sample the fitted chain to generate text (L25).

Every step is a module of this course applied to one model.

Learning outcomes

By the end of this lecture you will be able to:

1. estimate a transition matrix from counts,
2. derive ★ $\hat{P}_{ij} = \frac{n_{ij}}{n_{i.}}$ using Lagrange multipliers,
3. apply Dirichlet smoothing and distinguish prediction from MAP,
4. build a character n -gram model,
5. score held-out text in bits per character and perplexity,
6. sample new sequences from a fitted model,
7. identify how every course module appears in this one artifact,
8. and locate the next mathematical steps toward ML and DL.

**Estimate the transition matrix
from data**

An observed sequence

Two states A and B produced the sequence

A A B A B B B A

There are seven observed transitions:

$A \rightarrow A, A \rightarrow B, B \rightarrow A, A \rightarrow B,$

$B \rightarrow B, B \rightarrow B, B \rightarrow A.$

Count each ordered pair.

Transition count table

current	next A	next B	row total
A	1	2	3
B	2	2	4

Rows correspond to current states; columns correspond to next states, matching L25.

Normalize each row

For current state A:

$$\hat{P}_{AA} = \frac{1}{3}, \quad \hat{P}_{AB} = \frac{2}{3}.$$

For current state B:

$$\hat{P}_{BA} = \frac{2}{4} = \frac{1}{2}, \quad \hat{P}_{BB} = \frac{2}{4} = \frac{1}{2}.$$

Thus

$$\hat{P} = \begin{pmatrix} \frac{1}{3} & \frac{2}{3} \\ \frac{1}{2} & \frac{1}{2} \end{pmatrix}.$$

Why row normalization is plausible

Given the current state i , the next state must be one of the columns j .

The empirical conditional frequency is

$$\frac{\text{number of observed } i \rightarrow j \text{ transitions}}{\text{number of transitions leaving } i}.$$

The calculation is also the maximum-likelihood estimator (L14). We now derive it.

**Maximum likelihood for one
row**

Likelihood of transition counts

Let n_{ij} count transitions from state i to state j .

Ignoring the initial-state term, the path likelihood is

$$L(P) = \prod_i \prod_j P_{ij}^{n_{ij}}.$$

The log-likelihood is

$$\ell(P) = \sum_i \sum_j n_{ij} \log P_{ij}.$$

Constraints live row by row

Every row is a probability distribution:

$$\sum_j P_{ij} = 1, \quad P_{ij} \geq 0.$$

The log-likelihood separates across rows. For one fixed current state i , solve

$$\max_{P_{i1}, \dots, P_{iK}} \sum_j n_{ij} \log P_{ij}$$

$$\text{subject to } \sum_j P_{ij} = 1.$$

★ Form the row Lagrangian

$$\mathcal{L}_i = \sum_j n_{ij} \log P_{ij} - \lambda_i \left(\sum_j P_{ij} - 1 \right).$$

This is L19's convention for maximization, $\mathcal{L} = f - \lambda h$, with $h = \sum_j P_{ij} - 1$.

First assume every $n_{ij} > 0$, so the maximizing row is in the simplex interior.

Differentiate with respect to P_{ij} :

$$\partial \frac{\mathcal{L}_i}{\partial P_{ij}} = \frac{n_{ij}}{P_{ij}} - \lambda_i = 0.$$

Therefore

$$P_{ij} = \frac{n_{ij}}{\lambda_i}.$$

★ Use the row-sum constraint

D DERIVATION

$$\sum_j P_{ij} = \sum_j \frac{n_{ij}}{\lambda_i} = 1.$$

Define the row total

$$n_{i.} = \sum_j n_{ij}.$$

Then

$$\lambda_i = n_{i.}.$$

Substitute $\lambda_i = n_{i.}$:

$$\hat{P}_{ij} = \frac{n_{ij}}{n_{i.}} \quad \text{for } n_{i.} > 0.$$

The maximum-likelihood transition probability is the empirical fraction of departures from i that arrive at j .

If $n_{ij} = 0$, the KKT conditions (L20) place that coordinate on the boundary at $\hat{P}_{ij} = 0$, giving the same formula. If $n_{i.} = 0$, the row was never observed and its transition probabilities are not identifiable without a prior or another modeling choice.

Return to the A/B sequence

$$n_{A\cdot} = 3, \quad n_{B\cdot} = 4.$$

$$\hat{P}_{AA} = \frac{1}{3}, \quad \hat{P}_{AB} = \frac{2}{3},$$

$$\hat{P}_{BA} = \frac{2}{4}, \quad \hat{P}_{BB} = \frac{2}{4}.$$

The intuitive row-normalization calculation and the constrained MLE agree.

The initial distribution

The full sequence likelihood also contains $p(X_1)$.

Possible choices:

1. estimate an initial-state distribution from many independent sequences,
2. fix the observed first state,
3. initialize from a known external distribution,
4. use the stationary distribution.

Transition estimation and initialization answer different questions.

Code from a sequence

```
import numpy as np

path = np.array([0, 0, 1, 0, 1, 1, 1, 0])
counts = np.zeros((2, 2), dtype=int)

for current, nxt in zip(path[:-1], path[1:]):
    counts[current, nxt] += 1

P = counts / counts.sum(axis=1, keepdims=True)
print(counts)  # [[1 2], [2 2]]
print(P)       # [[.333 .667], [.5 .5]]
```

From state C, observed next-state counts are $(6, 3, 1)$. What is the MLE row?

A. $(6, 3, 1)$

B. $(0.6, 0.3, 0.1)$

C. $(\frac{1}{3}, \frac{1}{3}, \frac{1}{3})$

D. $(\frac{7}{13}, \frac{4}{13}, \frac{2}{13})$

Correct: B — $(0.6, 0.3, 0.1)$

Why: Divide each count by the row total $6 + 3 + 1 = 10$.

A bigram language model

Characters are states


Let each character be a state: space plus a through z.

A character bigram model assumes


$$p(x_t | x_1, \dots, x_{t-1}) \approx p(x_t | x_{t-1}).$$

The transition matrix has 27×27 entries. Row t stores probabilities for the character following t .

bigram counts in the course lecture plan (log colour scale)



A subset of the alphabet. Common spelling and spacing patterns appear as large counts.



On this corpus t is followed most often by i (as in *transition*, *optimization*), then by a space and by h .

Build character counts

```
from collections import defaultdict, Counter

text = clean(open("lecture-plan-detailed.md").read())
counts = defaultdict(Counter)

for current, nxt in zip(text[:-1], text[1:]):
    counts[current][nxt] += 1

row_t = counts["t"]
total_t = sum(row_t.values())
p_t = {ch: n/total_t for ch, n in row_t.items()}
```

Bigram likelihood

For text $x_1, \dots, x_T$,

$$p(x_1, \dots, x_T) = p(x_1) \prod_{t=2}^T P_{x_{t-1}, x_t}.$$

The training MLE simply counts every adjacent pair and normalizes rows.

No gradient descent is needed because the constrained optimum has a closed form.

Higher-order n -grams

1. unigram: $p(x_t)$ — no context,
2. bigram: $p(x_t|x_{t-1})$ — one previous character,
3. trigram: $p(x_t|x_{t-2}, x_{t-1})$ — two previous characters,
4. 4-gram: three previous characters.

A higher-order model is still first-order Markov after treating the whole context as the state.

Parameter growth

With alphabet size K and context length m , a dense table contains

$$K^{m(K-1)}$$

free transition parameters.

For $K = 27$:

1. bigram ($m = 1$): 702 free parameters,
2. trigram ($m = 2$): 18,954,
3. 4-gram ($m = 3$): 511,758.

Longer context needs much more data.

**Unseen does not mean
impossible**

The zero-probability problem

Suppose one row has counts

$(8, 2, 0, 0)$.


MLE gives

$(0.8, 0.2, 0, 0)$.

If either unseen transition occurs in test data, its log loss is infinite:

$$-\log 0 = \infty.$$

counts (8, 2, 0, 0): smoothing reserves mass for unseen outcomes



Add-one prediction changes (8, 2, 0, 0) to (9/14, 3/14, 1/14, 1/14).

Dirichlet prior

For a transition row P_i , choose

$$P_i \sim \text{Dirichlet}(\alpha_1, \dots, \alpha_K).$$

After counts $n_{i1}, \dots, n_{iK}$, conjugacy gives

$$P_i | \text{data} \sim \text{Dirichlet}(n_{i1} + \alpha_1, \dots, n_{iK} + \alpha_K).$$

This is the multinomial analogue of L15's Beta–Bernoulli update.

Posterior predictive probabilities

The probability of the next transition, averaged over the posterior, is

$$P(X_{\text{next}} = j | X_{\text{current}} = i, \text{data}) = \frac{n_{ij} + \alpha_j}{n_{i\cdot} + \sum_k \alpha_k}.$$

For symmetric $\alpha_j = \alpha$, this is add- α smoothing.

Predictive mean is not always MAP

For $\alpha_j > 1$, the row-wise MAP is

$$\frac{n_{ij} + \alpha_j - 1}{n_{i\cdot} + \sum_k \alpha_k - K}.$$

The posterior predictive mean uses $n_{ij} + \alpha_j$ instead.

“Add-one smoothing” usually refers to predictive probabilities under a uniform Dirichlet prior, not the MAP for that prior.

Choose the smoothing strength

Large α pulls rows strongly toward the prior mean.

Small α stays closer to counts but may assign extremely small probabilities to unseen transitions.

Choose α on validation data using held-out log loss, not training likelihood alone.

**Test on text not used for
counting**

Held-out cross-entropy

For a test sequence of length T and context length m , bits per character is

$$\text{BPC} = -\frac{1}{T - m} \sum_{t=m+1}^T \log_2 q(x_t | x_{t-m}, \dots, x_{t-1}).$$

This is L24's cross-entropy in bits, measured on held-out text. Lower BPC means the model assigns higher probability to the observed continuation.

Character-level perplexity follows L24's convention for bits:

$$\text{perplexity} = 2^{\text{BPC}}.$$

Score in log-space

Our test segment has 5,390 characters, hence 5,389 scored transitions.

The bigram model's product of their conditional probabilities is about

$$2^{-3.581 \times 5389} \approx 10^{-5810},$$

far below the smallest positive normal float64, about 10^{-308} (L2).

Sum log-probabilities; never form the probability product (L2).


Train/test split

For the course lecture plan:

1. clean to lowercase letters and spaces,
2. first 80%: estimate counts,
3. final 20%: evaluate,
4. use add-0.1 smoothing,
5. never update counts on the test segment.

Chronological splitting better reflects sequence prediction than random character shuffling.

Context reduces held-out coding cost



Held-out BPC falls from 4.18 for a unigram to 2.69 for a 4-gram on this corpus.

Convert BPC to effective choices

model	BPC	perplexity
unigram	4.181	18.1
bigram	3.581	12.0
trigram	2.991	7.9
4-gram	2.687	6.4

The longer context narrows the effective next-character choice set.

More context is not automatically better

A higher-order model can have lower training loss but worse test loss because many contexts are rare.

Control the trade-off with:

1. smoothing,
2. backoff or interpolation with shorter contexts,
3. validation data,
4. parameter sharing.

Neural language models replace explicit tables with learned shared representations.

Which comparison makes character-level BPC meaningful?

A. different test corpora

B. same corpus and same character alphabet

C. training loss for one model and test loss for another

D. different log bases without conversion

Correct: B — same corpus and same character alphabet

Why: Cross-entropy is meaningful only when the evaluated data and unit of a symbol are held fixed.

Sample the fitted model

Generation is repeated conditional sampling

Given current context c_t :

1. retrieve the smoothed row $q(\cdot | c_t)$,
2. sample the next character x_{t+1} ,
3. append it to the output,
4. update the context,
5. repeat.

Generation reads the same fitted, smoothed rows that scoring reads. No second model is trained.

Bigram sampler

```
rng = np.random.default_rng(8)
start, order = "the ", 1
context, output = start[-order:], list(start)

for _ in range(320):
    probs = smoothed_row(context, alpha=0.001)
    nxt = rng.choice(list(alphabet), p=probs)
    output.append(nxt)
    context = (context + nxt)[-order:]

print("".join(output))
```

Smoothing strength is a per-use choice: scoring used $\alpha = 0.1$ to keep held-out zeros finite; here $\alpha = 0.001$, because heavier smoothing would inject uniformly random characters into samples.

What a bigram learns

The sampler above, run on the course-plan bigram counts, prints

```
the fulothecanaie s denivala chenen ch gers vex herde inghar  
ne d heg pe tion ceadeximaty de itangsinim sata lly bin...
```

Local letter pairs look plausible; words and sentences do not. One previous character is insufficient for language structure.

A 4-gram sample

The same sampler with `order = 3` (three previous characters) prints

```
the exchang compostep newton derivation lives x ent inputed  
bound float does ch dernouilly ipynb invert convectors manythis...
```

Short fragments from the course vocabulary appear, but long-range syntax remains weak.

Swap in a larger corpus

Replace the course plan with Tiny Shakespeare (1.1 MB, 65 distinct characters, speaker names and line breaks kept):

```
text = open("input.txt", encoding="utf-8").read()
table = train_counts(text, order=5)
print(generate(table, order=5, n=200, seed=4,
               alpha=0.001, start=text[:5]))
```

The algorithm does not change. Corpus size, alphabet, and context order change the fitted probabilities.

Shakespeare from five-character contexts

The run above prints, verbatim:

```
First, yet this, for your wife, we met, were spake it.
```

```
AUTOLYCUS:
```

```
Very true subjects the oracle!
```

```
DUKE VINCENTIO:
```

```
Good made, but powerful old Grey his since he will not have order,  
Tullus not anon, we'll
```

Real speaker names, mostly real words, weak long-range syntax — the texture of Lecture 1's sample. The promise from L1 is kept: this model was trained from scratch, today, by counting.

What scale adds

Our best count model reaches 2.69 BPC on the held-out course plan.

Published character-level results for GPT-2 (1.5B parameters, zero-shot; Radford et al., 2019): 0.98 BPC on text8 and 0.93 bits per byte on enwik8.

Different corpora and alphabets, so the numbers are not directly comparable — but the objective is identical: minimize held-out cross-entropy (L24).

What changes between the two models is the parameterization of $q(x_t | \text{context})$ and the amount of data — not the loss, the evaluation, or the sampling procedure.

Every module appears here

Machine numbers and linear algebra

lecture	object	use in the sequence model
L2	floating point	sum log probabilities; never multiply thousands of tiny terms
L3	vectors	a state distribution is a probability vector
L4	matrices	$\mu_{t+1} = \mu_t P$
L5	eigenvectors	stationary distribution has eigenvalue one
L6	low rank	large transition tables may admit compressed structure

Calculus and automatic differentiation

lecture	object	use beyond the count model
L7	local linearization	optimize parameterized sequence models
L8	gradients	direction of cross-entropy change
L9	Jacobians/Hessians	sensitivity and curvature of model outputs
L10	finite/forward diff	check small derivatives
L11	reverse mode	backpropagate sequence loss through shared parameters

Probability and estimation

lecture	object	use in the sequence model
L12–13	distributions	represent uncertainty and dependence
L14	MLE	row-normalized transition counts
L15	MAP / conjugacy	Dirichlet smoothing of each row

Counting, normalization, and smoothing are estimation procedures—not ad hoc text tricks.

Optimization

lecture	object	use in the sequence model
L16	convexity	row log-likelihood has one global optimum
L17–18	first/second order	fit models without a count closed form
L19	Lagrange multipliers	derive row-normalized MLE
L20	KKT	handle probability non-negativity and active zeros
L21	structured solvers	recognize special training objectives

Information and sequences

lecture	object	use in the sequence model
L22	entropy	irreducible uncertainty of a source
L23	coding	probability becomes description length
L24	cross-entropy / KL	held-out BPC and model mismatch
L25	Markov chains	transition matrix, stationary distribution, PageRank
L26	learning sequences	estimate, smooth, score, and sample

From n -grams to neural language models

An n -gram model stores a separate probability row for every discrete context.

A neural language model:

1. maps tokens and contexts to vectors,
2. shares parameters across related contexts,
3. produces logits and softmax probabilities,
4. minimizes the same cross-entropy,
5. samples from the same conditional factorization.

The probability and loss framework remains; the parameterization becomes much richer.

**What you can now build and
check**

A compact sequence-model workflow

1. Define states or symbols and context length.
2. Split training, validation, and test sequences.
3. Count transitions or choose a parameterized model.
4. Normalize with appropriate smoothing.
5. Evaluate held-out cross-entropy in a clear unit.
6. Compare against simpler baselines.
7. Inspect generated samples and failure modes.
8. Record assumptions about stationarity and memory.

Transition-count MLE is row normalization; smoothing makes prediction finite; cross-entropy evaluates the fitted sequence model.

1. For an observed row, $\hat{P}_{ij} = \frac{n_{ij}}{n_{i\cdot}}$ follows from constrained maximum likelihood.
2. A Dirichlet prior yields add- α posterior predictive probabilities.
3. Longer context can reduce BPC but increases parameter and data requirements.
4. Generation repeatedly samples the same conditional probabilities used for scoring.
5. The character model uses every mathematical module in the course.

Reading — MacKay Ch. 2, revisited now that every tool in it is familiar. Optional: Karpathy, *makemore* part 1 (video) — the same bigram model, first by counting, then by gradient descent.

Practice

1. Estimate P from transitions $A \rightarrow A$ twice, $A \rightarrow B$ once, $B \rightarrow A$ three times, and $B \rightarrow B$ once.
2. Apply add-one smoothing to the A row.
3. A test sequence receives probabilities $(\frac{1}{2}, \frac{1}{4}, \frac{1}{8})$. Find total bits and average BPC.
4. Why can a 10-gram have worse held-out BPC than a 4-gram?

1. $\hat{P} = \begin{pmatrix} \frac{2}{3} & \frac{1}{3} \\ \frac{3}{4} & \frac{1}{4} \end{pmatrix}$.
2. Counts (2, 1) become predictive probabilities $(\frac{3}{5}, \frac{2}{5})$.
3. Total = 1 + 2 + 3 = 6 bits; average = 2 BPC.
4. Most long contexts are rare or unseen, so their estimates have high variance unless data sharing, smoothing, or backoff is strong enough.

What follows this course

In **ES 335 (Machine Learning)**, these foundations support:

1. linear and logistic regression,
2. regularization and model selection,
3. support vector machines and kernels,
4. tree ensembles and probabilistic models.

In **ES 667 (Deep Learning)**, they support:

1. neural networks and backpropagation,
2. embeddings and attention,
3. convolutional and sequence models,
4. generative models and large language models.

Attention, seen from this course

Our n -gram fixes the context: the previous m symbols, no more.

Attention is a smarter memory than the one-step Markov context — the model learns which parts of the history to condition on.

The output of that memory still feeds the factorization, loss, and sampler you used today. ES 667 builds the memory.

Final checklist

You can now ask of an AI method:

1. What is represented as a vector or matrix?
2. Which probability model produced the loss?
3. Which derivative is being computed, and in what shape?
4. What local approximation justifies the update?
5. Which constraint or numerical scale can break it?
6. What does its cross-entropy say in bits?
7. Which assumptions connect observations across time?

Given probabilities q_j , temperature $\tau > 0$ defines

$$q_j^\tau \propto \exp\left(\frac{\log q_j}{\tau}\right).$$

1. $\tau < 1$: sharper, more repetitive samples,
2. $\tau = 1$: original model,
3. $\tau > 1$: flatter, more varied, more errors.

Temperature changes generation; it does not improve held-out likelihood.

A hidden Markov model keeps the Markov chain but hides it: the chain runs over states Z_t , and we observe only emissions $X_t \sim p(x|Z_t)$.

Three standard questions — score a sequence, find the best hidden path, fit the parameters — are all answered with this course's tools: conditional probability and dynamic programming over log-probabilities.

Reference: Rabiner (1989), *A tutorial on hidden Markov models*. HMMs are covered properly in probabilistic-ML courses; this slide is a pointer.

The course in 26 lines

- L1** Every AI system is a stack of math — by L26 you'll have built one from scratch.
- L2** Computers don't do real numbers — they do ~4 billion of them, unevenly spaced.
- L3** Similarity is an angle.
- L4** A matrix is a verb, not a table.
- L5** Eigenvectors are the directions a matrix can't turn.
- L6** Every matrix is a rotation, a stretch, and a rotation.
- L7** Calculus is the art of replacing a function with a line.
- L8** The gradient points where the contour lines crowd.
- L9** Curvature is a matrix, and its sign is the shape of the bowl.
- L10** Finite differences balance truncation against floating-point rounding.
- L11** Backprop is just the chain rule with memoization.
- L12** Continuous probability is area, and transformations charge a Jacobian toll.
- L13** A Gaussian is a point and an ellipse.
- L14** Many standard supervised losses are negative log-likelihoods for explicit observation models.
- L15** L2 and L1 penalties are negative log-priors in MAP estimation.
- L16** For a convex objective, every local minimum is global.
- L17** The learning rate is how far you trust a linear approximation.
- L18** Newton doesn't step downhill — it teleports to the bottom of its local bowl.
- L19** At the optimum, the objective and the constraint pull in the same direction.
- L20** KKT is Lagrange with inequalities and receipts.
- L21** Don't solve optimization problems — recognize them.
- L22** Information is surprise; entropy is its average.
- L23** A good predictor and a good compressor are the same object.
- L24** Cross-entropy is the price of believing the wrong distribution.
- L25** A Markov chain is a matrix; its destiny is an eigenvector.
- L26** You just trained a language model — everything else is scale.

You just trained a language model — everything else is scale.

Next: **Tutorial 13** — fit, score, and sample this model yourself. Then the endsem — and ES 335, whenever you are ready.